

Unit 2

Process Management and Scheduling

Contents


1. Process Concepts
2. Process Scheduling
3. Operations on Process
4. Inter- Process Communication
5. Threads
6. Multi Threading Models
7. Basic Concepts on Process Scheduling
8. Scheduling Criteria
9. Scheduling Algorithms
10. Algorithm Evaluation

2.1 Process Concept

- Process - An instance of **computer program in execution**
- Some times also called as jobs or task
 - Batch system – **jobs**
 - Time-shared systems – **user programs or tasks**
- Process on modern systems can have several threads of execution or several different pieces of the program running in parallel
- Each process has its own private resources, data and associated statistics
- A OS creates a data structure in the memory

Process Concepts – The Process


- In memory, a process consists of **multiple parts**:
 - **Program code**, also called **text section**
 - **Stack** containing temporary data
 - Function parameters, return addresses, local variables
 - **Data section** containing global variables
 - **Heap** containing memory dynamically allocated during run time



Process Concept - Process State

- As a process executes, it changes *state*
 - **new**: The process is being created
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **ready**: The process is waiting to be assigned to a processor
 - **terminated**: The process has finished execution

Diagram of Process State




Process Concepts - Process Control

Block (PCB)

Information associated with each process

- **Process state** – new, ready, running...
- **Process Number** – process id, parent process
- **Program counter** – next instr. for exec
- **CPU registers** – no. and type of reg. used
- **Memory-management information**
 - Limit of page table and segment table
- **Accounting information**
 - Amt of CPU time used, time limit
- **I/O status information**
 - List of I/O devices



Process Concepts - Threads

- Threads – lightweight process
- So far, process has a single thread of execution
 - One task at a time
- Consider having multiple program counters per process
 - Multiple locations can execute at once
 - Multiple tasks at a time
- PCB must be extended to handle threads:
 - Store thread details
 - Multiple program counters


2.2 Process Scheduling

- Maximize CPU use, quickly switch processes onto CPU for time sharing
- **Process scheduler** selects among available processes for next execution on CPU
 - Scheduling Queues
 - Schedulers
 - Context Switch


Process Scheduling – Scheduling Queues

- Process Maintains **scheduling queues** of processes
 - **Job queue** – set of all processes in the system
 - **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
 - **Device queues** – set of processes waiting for an I/O device
 - Processes migrate among the various queues
- Queue is generally stored as linked list

Ready Queue And Various I/O Device Queues



Representation of Process Scheduling



Process Scheduling- Schedulers


- **Long-term scheduler** (or job scheduler) – selects which processes should be brought into the ready queue
- **Short-term scheduler** (or CPU scheduler) – selects which process should be executed next and allocates CPU
 - Sometimes the only scheduler in a system
- The Schedulers will implement scheduling policies.

Schedulers (Cont.)

- **Short-term scheduler** is invoked very frequently (milliseconds)
⇒ (must be fast)
- **Long-term scheduler** is invoked very infrequently (seconds, minutes) ⇒ (may be slow)
- The long-term scheduler controls the *degree of multiprogramming*
- Processes can be described as either:
 - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process** – spends more time doing computations; few very long CPU bursts


Addition of Medium Term Scheduling

- **Medium-term scheduler** (or job scheduler) – removes process from memory so that reduce the degree of multiprogramming. Later, it reintroduces in main memory. i.e, it takes care of swapping process



Process Scheduling - Context Switch

- When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process via a **context switch**.
- **Context** of a process represented in the PCB
- Context-switch time is overhead; the system does no useful work while switching
 - The more complex the OS and the PCB -> longer the context switch




2.3 Operations on Processes

- A process in most of the systems can execute concurrently, and they must be created & deleted dynamically. Thus, these Systems must provide mechanism for
 - Process Creation
 - Process Deletion

Operations on Processes – Process Creation


- A (**parent**) process can create several (**children**) processes
 - Children can, in turn, create other processes
 - Hence, a **tree** of processes forms
- Generally, process identified and managed via a **process identifier (pid)**
- Resource sharing options (of process creation)
 - Parent and children share all resources
 - Children share subset of parent's resources
 - One usage is to prevent system overload by too many child processes
 - Parent and child share no resources
- Execution options
 - Parent and children execute concurrently
 - Parent waits until children terminate

Operations on Processes – Process Creation



Operations on Processes – Process Creation

- Address space options
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - **fork()** system call creates new process
 - Child is a copy of parent's address space
 - except fork() returns 0 to child and nonzero to parent
 - **exec()** system call used after a **fork()** to replace the process' memory space with a new program



Operations on Processes – Process Termination

- Process executes last statement and then asks the operating system to delete it using the **exit()** system call.
 - Returns status data from child to parent (via **wait()**)
 - Process' resources are deallocated by operating system
- Parent may terminate the execution of children processes using the **abort()** system call. Some reasons for doing so:
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - The parent is exiting and the operating systems does not allow a child to continue if its parent terminates

Operations on Processes – Process Termination


- Some OS's don't allow child to exists if its parent has terminated
 - **cascading termination** - if a process terminates, then all its children, grandchildren, etc must also be terminated.
 - The termination is initiated by the operating system
- The parent process **may** wait for termination of a child process by using the **wait()** system call .
 - The call returns status information and the pid of the terminated process
pid = wait(&status);
- Process finished its execution but the state not changed is a **zombie**
- If parent terminated without invoking **wait**, process is an **orphan**

2.4 Inter Process Communication

- Processes within a system may be *independent* or *cooperating*
 - When processes execute they produce some **computational results**
 - *Independent* process cannot affect (or be affected) by results of another process
 - *Cooperating* process can affect (or be affected) by results of another process
- Advantages of process cooperation
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience

-contd..

- For fast exchange of information, cooperating processes need some **interprocess communication (IPC) mechanisms**
- Two models of IPC
 - Shared memory**
 - Message passing**



IPC – Shared Memory Systems

- An area of memory is shared among the processes that wish to communicate
- The communication is under the control of **the users processes**, not the OS.
- Major issue is to provide mechanism that will allow the user processes to **synchronize** their actions when they access shared memory.

IPC – Shared Memory Systems

- **Producer-consumer problem** – a common paradigm for cooperating processes
- **Producer** process : produces some information
- **Consumer** process : consumes this information
- **Challenge:**
 - Producer and consumer should run **concurrently** and **efficiently**
 - Producer and consumer must be **synchronized**
 - Consumer cannot consume an item before it is produced
- Shared-memory is solution to producer-consumer
- Uses a buffer in shared memory to exchange information
- **unbounded-buffer** assumes no practical limit on the buffer size
- **bounded-buffer** assumes a fixed buffer size

IPC – Message Passing

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - send(message)
 - receive(message)
- If P and Q wish to communicate, they need to:
 - Establish a communication link between them
 - Exchange messages via send/receive

IPC – Message Passing

- Implementation of Communication link
 - Physical link (eg: Shared memory, hardware bus)
 - Logical link(eg: direct or indirect, synchronous or asynchronous)
- The message size is either **fixed or variable**
- Generally variable sized messages are easy to implement
- Implementation issues:
 - How are links established?
 - Can a link be associated with more than two processes?
 - How many links can there be between every pair of communicating processes?
 - What is the capacity (buffer size) of a link?
 - Is the size of a message that the link can accommodate fixed or variable?
 - Is a link unidirectional or bi-directional?

IPC MessagePassing – Naming- Direct Communication

- Processes must name each other explicitly:
 - **send** ($P, message$) – send a message to process P
 - **receive**($Q, message$) – receive a message from process Q
- Properties of a direct communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link is usually bi-directional

IPC MessagePassing- Naming – Indirect Communication

- Messages are directed and received from **mailboxes** (also referred to as **ports**)
 - Each mailbox has a **unique id**
 - Processes can communicate only if they share a mailbox
- Properties of an indirect communication link
 - Link established only if both members of the pair have a shared mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional

IPC – Message Passing- Indirect Communication

Operations

create a new mailbox (port)

send and receive messages through mailbox

destroy a mailbox

Primitives are defined as:

send(*A, message*) – send a message to mailbox A

receive(*A, message*) – receive a message from mailbox A

IPC – Message Passing - Synchronization

- Message passing may be either
 - **Blocking**, or
 - **Non-blocking**
- **Blocking** is considered **synchronous**
 - **Blocking send** -- the sender is blocked until the message is received
 - **Blocking receive** -- the receiver is blocked until a message is available
- **Non-blocking** is considered **asynchronous**
 - **Non-blocking send** -- the sender sends the message and continues sending
 - **Non-blocking receive** -- the receiver receives:
 - A valid message, or
 - Null message
- Different combinations possible
 - If both send and receive are blocking – called a **rendezvous**

IPC – Message Passing - Buffering

- Queue of messages is attached to the link.
- Implemented in one of three ways
 1. **Zero capacity** – no messages are queued on a link.
 - Sender must wait for receiver
 2. **Bounded capacity** – finite length of n messages
 - Sender must wait if link full
 3. **Unbounded capacity** – infinite length
 - Sender never waits

Basics Concepts in Process Scheduling

- Maximum CPU utilization obtained with multiprogramming

- waiting for I/O is wasteful
 - 1 thread will utilize only 1 core

- CPU–I/O Burst Cycle

- Process execution consists of:
 - a **cycle** of CPU execution
 - and I/O wait

- **CPU burst** followed by **I/O burst**

- CPU burst distribution is of main concern

**load store
add store
read from file**

wait for I/O

**store increment
index
write to file**

wait for I/O

**load store
add store
read from file**

wait for I/O

CPU burst

I/O burst

CPU burst

I/O burst

CPU burst

I/O burst

Basics Concepts in Process Scheduling

CPU Scheduler

- **Short-term scheduler**
 - Selects 1 process from the ready queue
 - then allocates the CPU to it
 - Queue may be ordered in various ways : FIFO Queue, Priority Queue, tree, linked list
- CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from waiting to ready
 4. Terminates
- Scheduling under 1 and 4 is called **nonpreemptive (=cooperative)**
- All other scheduling is called **preemptive**
 - Process can be **interrupted** and must release the CPU
 - Special care should be taken to prevent problems that can arise
 - Access to shared data – race condition can happen, if not handled

Basics Concepts in Process Scheduling

Dispatcher

- **Dispatcher**
 - a module that gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - switching context
 - switching to user mode
 - jumping to the proper location in the user program to restart that program
- **Dispatch latency**
 - Time it takes for the dispatcher to stop one process and start another running
 - This time should be as small as possible

Scheduling Criteria

- How do we decide which scheduling algorithm is good?
- Many **criteria** for judging this has been suggested
 - Which characteristics considered can **change significantly** which algorithm is considered the best
- **CPU utilization** – keep the CPU as busy as possible
- **Throughput** – # of processes that complete their execution per time unit
- **Turnaround time** – amount of time to execute a particular process
- **Waiting time** – amount of time a process has been waiting in the ready queue
- **Response time** – amount of time it takes to start responding
 - Used for interactive systems
 - Time from when a request was submitted until the first response is produced

Scheduling Criteria

Scheduling Algorithm Optimization Criteria

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time

Scheduling Algorithm

First- Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- Suppose that the processes arrive in the order: P_1, P_2, P_3
The [Gantt Chart](#) for the schedule is:



- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$

Scheduling Algorithm FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

$$P_2, P_3, P_1$$

- The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6; P_2 = 0; P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case
 - Hence, average waiting time of FCFS not minimal
 - And it may vary substantially
- FCFS is **nonpreemptive**
 - Not a good idea for timesharing systems
- FCFS suffers from **the convoy effect**, shorter job to wait for a longer time

Scheduling Algorithm Shortest-Job-First (SJF)

Scheduling

- Associate with each process the length of its **next CPU burst**
 - SJF uses these lengths to schedule the process with the shortest time
- Notice, the **burst** is used by SJF,
 - **not** the process end-to-end running time
 - implied by word “job” in SJF
 - Hence, it should be called ‘‘Shorted-Next-CPU-Burst’’
 - However, “job” is used for historic reasons
- Two versions of SJF: **preemptive** and **nonpreemptive**
 - Assume
 - A new process Pnew arrives while the current one Pcur is still executing
 - The burst of Pnew is less than what is left of Pcur
 - Nonpreemptive SJF – will let Pcur finish
 - Preemptive SJF – will preempt Pcur and let Pnew execute
 - This is also called **shortest-remaining-time-first** scheduling


Scheduling Algorithm SJF (Cont.)

- **Advantage:**
 - SJF is **optimal** in terms of the average waiting time
- **Challenge of SJF:**
 - Hinges on knowing the length of the next CPU burst
 - But how can we know it?
 - Solutions: ask user or estimate it
 - In a **batch system** and **long-term scheduler**
 - Could ask the user for the job time limit
 - The user is motivated to accurately estimate it
 - Lower value means faster response
 - Too low a value will cause time-limit violation and job rescheduling
 - In a **short-term scheduling**
 - Use **estimation**
 - Will be explained shortly

Scheduling Algorithm Example of SJF

<u>Process</u>	<u>Burst Time</u>
P_1	6
P_2	8
P_3	7
P_4	3

- SJF scheduling chart



- Average waiting time = $(3 + 16 + 9 + 0) / 4 = 7$


Scheduling Algorithm Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
 - Preemptive
 - Nonpreemptive
- SJF is priority scheduling where priority is the inverse of predicted next CPU burst time
- Problem \equiv **Starvation** – low priority processes may never execute
- Solution \equiv **Aging** – as time progresses increase the priority of the process

Scheduling Algorithm

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

- Priority scheduling Gantt Chart



- Average waiting time = 8.2 msec

Scheduling Algorithm Round Robin (RR)


- Each process gets a small unit of CPU time
 - **Time quantum** q
 - Usually 10-100 milliseconds
- After this time has elapsed:
 - the process is **preempted** and
 - added to the end of the ready queue
- The process might run for $\leq q$ time
 - For example, when it does I/O
- **If**
 - n processes in the ready queue, and
 - the time quantum is q
- **then**
 - “Each process gets $1/n$ of the CPU time”
 - **Incorrect statement from the textbook**
 - in chunks of $\leq q$ time units at once
 - Each process waits $\leq (n-1)q$ time units

Scheduling Algorithm

Example of RR with Time Quantum = 4

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- The Gantt chart is:



- Typically:
 - Higher **average turnaround** (end-to-end running time) than SJF
 - But better **response** than SJF

Scheduling Algorithm - Multilevel Queue

- Another class of scheduling algorithms when processes are classified into groups, for example:
 - **foreground** (interactive) processes
 - **background** (batch) processes
- Ready queue is partitioned into separate queues, e.g.:
 - **Foreground** and **background** queues
- Process is permanently assigned to one queue
- Each queue has its own scheduling algorithm, e.g.:
 - foreground – RR
 - background – FCFS

Scheduling Algorithm - Multilevel Queue

Scheduling must be done between the queues:

Fixed priority scheduling

For example, foreground queue might have absolute priority over background queue

Serve all from foreground then from background

May lead to starvation

Time slice scheduling


Each queue gets a certain amount of CPU time which it can schedule amongst its processes,
e.g.:

80% to foreground in RR

20% to background in FCFS

Scheduling Algorithm - Multilevel Queue

highest priority



lowest priority

Scheduling Algorithm - Multilevel Feedback Queue

- The previous setup: a process is permanently assigned to one queue
 - **Advantage:** Low scheduling overhead
 - **Disadvantage:** Inflexible
- **Multilevel-feedback-queue** scheduling algorithm
 - Allows a process to move between the various queues
 - More flexible
 - **Idea:** separate processes based on the characteristics of their CPU bursts
 - If a process uses too much CPU time => moved to lower-priority queue
 - Keeps I/O-bound and interactive processes in the high-priority queue
 - A process that waits too long can be moved to a higher priority queue
 - This form of **aging** can prevent starvation

Scheduling Algorithm - Multilevel Feedback Queue

- **Multilevel-feedback-queue** scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - method used to determine which queue a process will enter when that process needs service
- **Multilevel-feedback-queue** scheduler
 - The most general CPU-scheduling algorithm
 - It can be configured to match a specific system under design
 - Unfortunately, it is also the most complex algorithm
 - Some means are needed to select values for all the parameters

Scheduling Evaluation

- How to select CPU-scheduling algorithm for an OS?
 - Determine **criteria**, then evaluate algorithms
- Evaluation Methods
 - Deterministic modeling
 - Queuing models
 - Simulations
 - Implementation